

1. **UNSOLVED TASKS:** The following tasks remained unsolved through up to 80 iterations of instruction changes, showing that the harness is not the source of error, but instead the tests contain contradictions or are missing information:
 - a. **teleport:** There is a direct contradiction here. The test is looking for ErrLimitReached for limit 5 on “hello”, but requirement 3 says to return for no error when the limit allows reading all content.
 - b. **ansible-83909:** In the requirement, the exact wording of the message is not given, and the wrong flag is provided. The test checks text and not behavior.
 - c. **qutebrowser-21b4:** The task does not specify an exact string in terms of variable names (values -> vmap) and data structure (calling .values()).
2. **DIFFERENCES:** What I did differently from the handout:
 - a. 0: interface and logging
 - i. Passes problem requirements and interface since they were double-encoded in one string
 - ii. Writes run_end and logs tool_call before the tool runs, otherwise there could be logs with no ending
 - iii. Improved git safety by refusing state-changing git commands to ensure patch is not lost
 - b. 1.1: tools
 - i. **read_file:** uses 400-line windows with a “continue” note instead of the whole file to avoid truncating middle of files in output
 - ii. **edit_file:** instead of an exact replace, preserves line endings so that entire patch rewrites are not necessary
 - iii. **repro_check:** implemented a suggested tool with check that should fail pre-fix and pass after
 - iv. **done** runs a final review and default-export check instead of just stopping
 - c. 1.2: the loop
 - i. Follow up reminders are sent instead of stopping the loop when there is no tool call, since the model stops early sometimes
 - d. 1.3: conversation state and prompt
 - i. Instead of just truncating, tool outputs with 80k+ tokens become stubs in a batch to keep the cache working
3. **SOLVED TASKS:** How I made the two solved_after_change.json tasks pass
 - a. Unfortunately, I was not able to find the two solved_after_change tasks. I used an earlier version of the harness that passed 4/10 tasks as the “before”, and the final harness that passed 7/10 tasks as the “after”. 30 candidates were tested across 16 repos, but none lead to a clear flip from fail to pass. Finally, my API budget ran out before I could test a reconstructed VO of the harness to test more pairings and identify what led to successful flips. I am so sorry about this!